


Advanced GPU Topics #1

Jeremy Appleyard, September 2015



In this talk

- Profiling GPU applications
- Optimizing Data Movement
- Using MPI with GPUs

Ask questions at any point!

Profiling GPU Applications

Profiling Tools

Many options!

From NVIDIA

- nvprof
- NVIDIA Visual profiler
 - Standalone (nvvp)
 - Integrated into Nsight Eclipse Edition (nsight)
- Nsight Visual Studio Edition

Third Party

- System
- VampirTrace
- PAPI CUDA component
- HPC Toolkit

This talk

- ▶ We will focus on nvprof and nvvp
- ▶ nvprof => NVIDIA profiler
 - ▶ Command line
- ▶ nvvp => NVIDIA Visual Profiler
 - ▶ GUI based

nvprof

Simple usage

- `nvprof ./<executable>`
- Example: vector addition from yesterday's talk

```
attributes(global) subroutine vecAdd_GPU(c, a, b, n)
  INTEGER, value :: n
  REAL, device, intent(in)  :: a(n), b(n)
  REAL, device, intent(out) :: c(n)

  INTEGER :: i

  i = (blockIdx%x - 1) * blockDim%x + threadIdx%x

  if (i <= n) then
    c(i) = a(i) + b(i)
  end if
end subroutine vecAdd_GPU
```

```
nvprof ./vecAdd
```

```
==34092== NVPROF is profiling process 34092, command: ./vecAdd
```

```
==34092== Profiling application: ./vecAdd
```

```
==34092== Profiling result:
```

Time(%)	Time	Calls	Avg	Min	Max	Name
69.07%	3.8547ms	2	1.9274ms	1.9093ms	1.9454ms	[CUDA memcpy HtoD]
29.04%	1.6205ms	1	1.6205ms	1.6205ms	1.6205ms	[CUDA memcpy DtoH]
1.19%	66.178us	1	66.178us	66.178us	66.178us	kernel_vecadd_gpu_
0.71%	39.650us	2	19.825us	19.617us	20.033us	__pgi_dev_cumemset_4

```
==34092== API calls:
```

Time(%)	Time	Calls	Avg	Min	Max	Name
95.29%	321.93ms	3	107.31ms	215.28us	321.48ms	cudaMalloc
2.70%	9.1355ms	3	3.0452ms	2.1550ms	3.5234ms	cudaMemcpy
1.47%	4.9710ms	498	9.9810us	177ns	498.93us	cuDeviceGetAttribute
0.22%	758.16us	3	252.72us	234.62us	284.11us	cudaFree
0.15%	519.49us	6	86.581us	82.801us	89.206us	cuDeviceTotalMem
0.13%	449.63us	6	74.938us	70.892us	81.049us	cuDeviceGetName
0.02%	69.908us	3	23.302us	11.485us	43.331us	cudaLaunch
0.00%	12.300us	10	1.2300us	179ns	8.0460us	cudaSetupArgument
0.00%	3.4430us	12	286ns	193ns	716ns	cuDeviceGet
0.00%	2.8280us	2	1.4140us	285ns	2.5430us	cuDeviceGetCount
0.00%	2.5750us	3	858ns	448ns	1.6120us	cudaConfigureCall

```

nvprof ./vecAdd
==34092== NVPROF is profiling process 34092, command: ./vecAdd
==34092== Profiling application: ./vecAdd
==34092== Profiling result:
Time(%)      Time      Calls      Avg      Min      Max  Name
69.07%    3.8547ms         2  1.9274ms  1.9093ms  1.9454ms  [CUDA memcpy HtoD]
29.04%    1.6205ms         1  1.6205ms  1.6205ms  1.6205ms  [CUDA memcpy DtoH]
 1.19%    66.178us         1   66.178us  66.178us  66.178us  kernel_vecadd_gpu_
 0.71%    39.650us         2   19.825us  19.617us  20.033us  __pgi_dev_cumemset_4

```

- Top half of the profile is runtime measured from the GPU perspective
- What can we see?
 - Memcpy XtoX is memory copy to and from the GPU
 - The vector addition is only 1.19%!
 - This is valuable information!


```

nvprof ./vecAdd
==34092== NVPROF is profiling process 34092, command: ./vecAdd
==34092== Profiling application: ./vecAdd
==34092== Profiling result:
Time(%)      Time       Calls          Avg          Min          Max   Name
 69.07%    3.8547ms         2    1.9274ms    1.9093ms    1.9454ms  [CUDA memcpy HtoD]
 29.04%    1.6205ms         1    1.6205ms    1.6205ms    1.6205ms  [CUDA memcpy DtoH]
  1.19%    66.178us         1    66.178us    66.178us    66.178us  kernel_vecadd_gpu_
  0.71%    39.650us         2    19.825us    19.617us    20.033us  __pgi_dev_cumemset_4

```

- (PGI) OpenACC kernels will be named after the subroutine name and line number
- **eg.** `vecadd_17_gpu`
 - Subroutine `vecadd`
 - Line 17 in the file
 - Compiled for the `gpu`

- Bottom half of the profile is runtime measured from the CPU perspective
- What can we see?
 - First allocation is expensive (cuda initialisation)

==34092== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
95.29%	321.93ms	3	107.31ms	215.28us	321.48ms	cudaMalloc
2.70%	9.1355ms	3	3.0452ms	2.1550ms	3.5234ms	cudaMemcpy
1.47%	4.9710ms	498	9.9810us	177ns	498.93us	cuDeviceGetAttribute
0.22%	758.16us	3	252.72us	234.62us	284.11us	cudaFree
0.15%	519.49us	6	86.581us	82.801us	89.206us	cuDeviceTotalMem
0.13%	449.63us	6	74.938us	70.892us	81.049us	cuDeviceGetName
0.02%	69.908us	3	23.302us	11.485us	43.331us	cudaLaunch
0.00%	12.300us	10	1.2300us	179ns	8.0460us	cudaSetupArgument
0.00%	3.4430us	12	286ns	193ns	716ns	cuDeviceGet
0.00%	2.8280us	2	1.4140us	285ns	2.5430us	cuDeviceGetCount
0.00%	2.5750us	3	858ns	448ns	1.6120us	cudaConfigureCall

nvprof

More advanced options

- `nvprof -h`
- There are quite a few options!
- Some useful ones:
 - `-o`: creates an output file which can be imported into nvvp
 - `-m` and `-e`: collect metrics or events
 - `--analysis-metrics`: collect all metrics for import into nvvp
 - `--query-metrics` and `--query-events`: query which metrics/events are available

nvprof

Events and Metrics

- ▶ Most are quite in-depth, however some useful ones for quick analysis
- ▶ In general, events are only for the expert. Rarely useful
- ▶ (A few) useful metrics:
 - ▶ `dram_read_throughput`: Main GPU memory read throughput
 - ▶ `dram_write_throughput`: Main GPU memory write throughput
 - ▶ `flop_count_sp`: Number of single precision floating point operations
 - ▶ `flop_count_dp` : Number of double precision floating point operations

nvprof

Metrics Example

```
nvprof --metrics dram_read_throughput,dram_write_throughput,flop_count_sp,flop_count_dp ./vecAdd
```

```
==32928== NVPROF is profiling process 32928, command: ./vecAdd
```

```
==32928== Warning: Some kernel(s) will be replayed on device 0 in order to collect all events/metrics.
```

```
==32928== Profiling application: ./vecAdd
```

```
==32928== Profiling result:
```

```
==32928== Metric result:
```

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "Tesla K40m (0)"					
	Kernel: kernel_vecadd_gpu_				
1	dram_read_throughput	Device Memory Read Throughput	136.55GB/s	136.55GB/s	136.55GB/s
1	dram_write_throughput	Device Memory Write Throughput	70.062GB/s	70.062GB/s	70.062GB/s
1	flop_count_sp	Floating Point Operations(Single Precisi	1000000	1000000	1000000
1	flop_count_dp	Floating Point Operations(Double Precisi	0	0	0

nvprof

Metrics Example

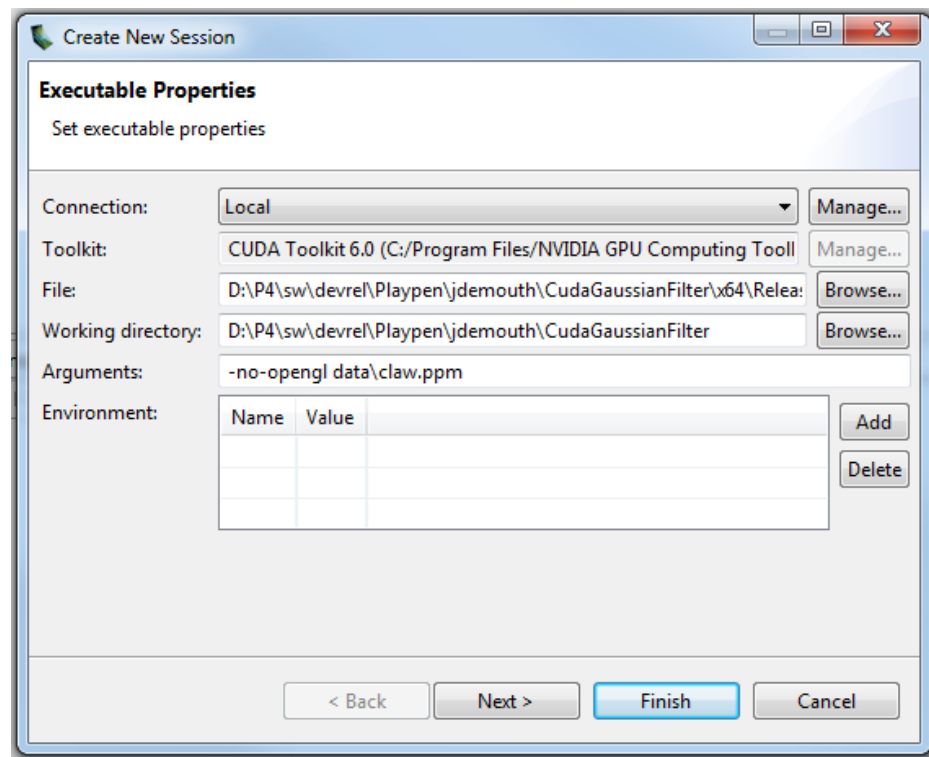
```
dram_read_throughput 136.55GB/s
dram_write_throughput 70.062GB/s
      flop_count_sp 1000000
      flop_count_dp 0
```

- ▶ Problem was addition of 1,000,000 element single precision vectors
 - ▶ 1,000,000 single precision flop count is therefore expected!
 - ▶ ~50 GFLOP/s
 - ▶ About 1% peak FLOPs
- ▶ Measured dram throughput is ~206GB/s
 - ▶ Peak for this machine is 288 GB/s
 - ▶ 72% peak bandwidth

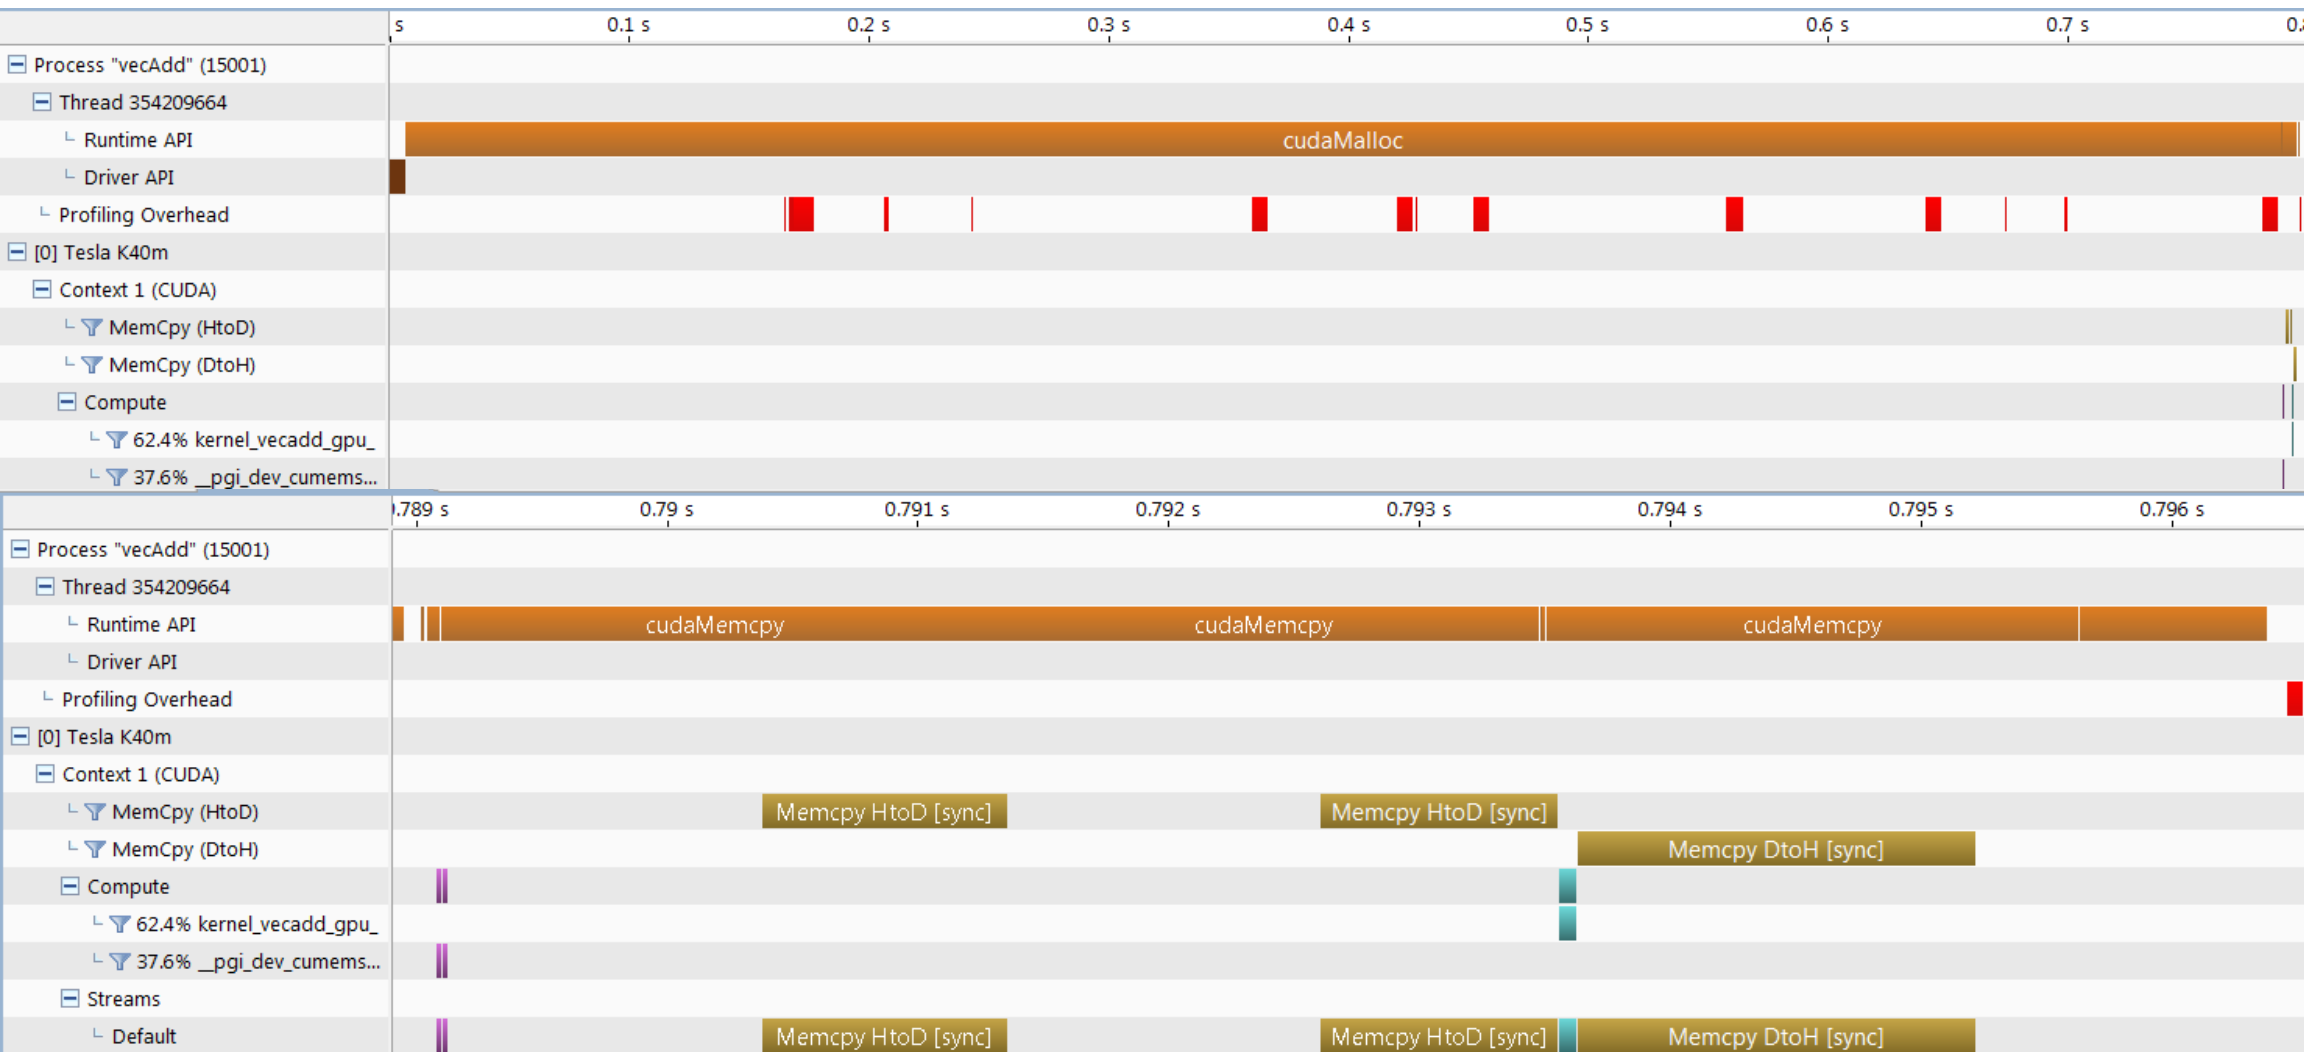
Visual Profiler

nvvp


- Either import the output of `nvprof -o ...`
 - Need to select metrics in advance
- Or run application through GUI
 - Re-run on-the-fly to compute new metrics based on requirements



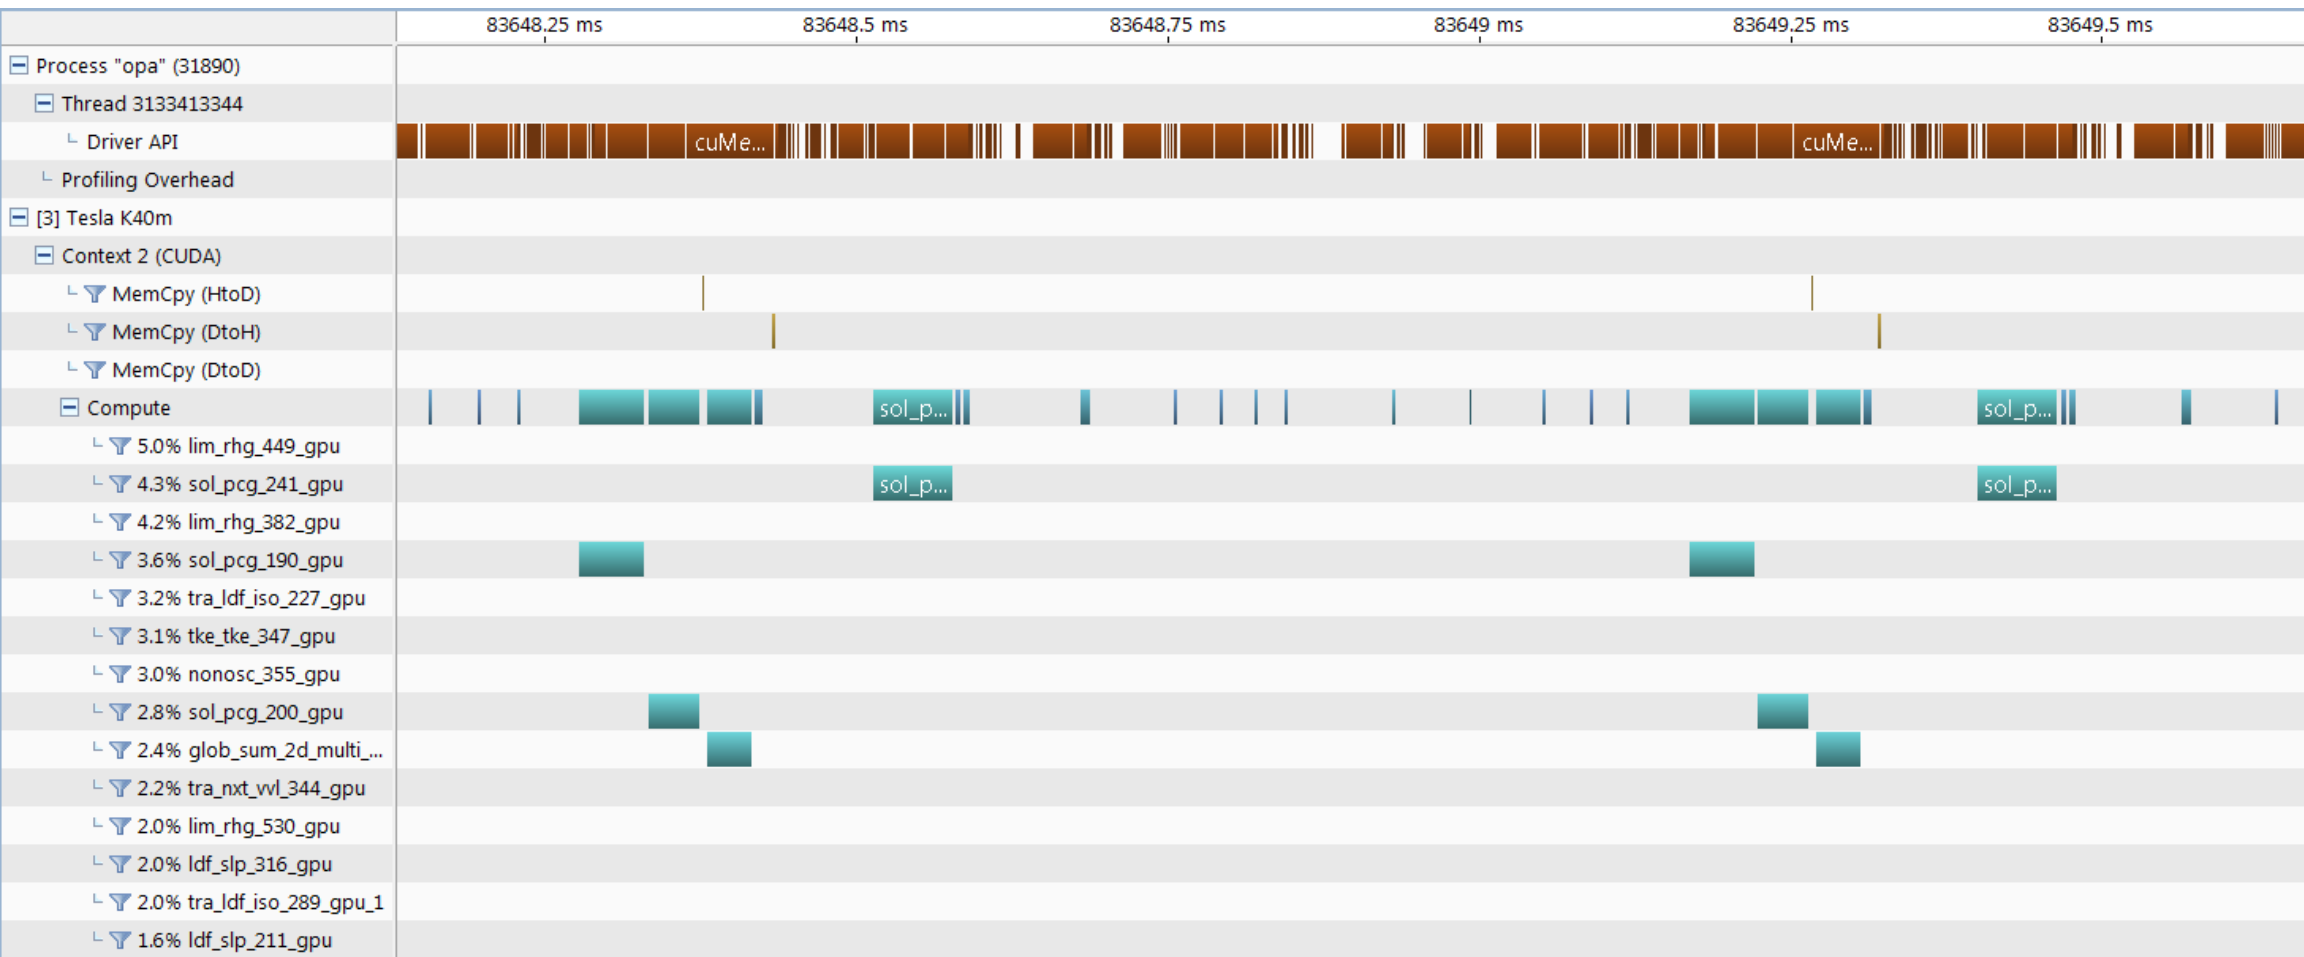
Vecadd - timeline



NEMO - timeline



NEMO - details

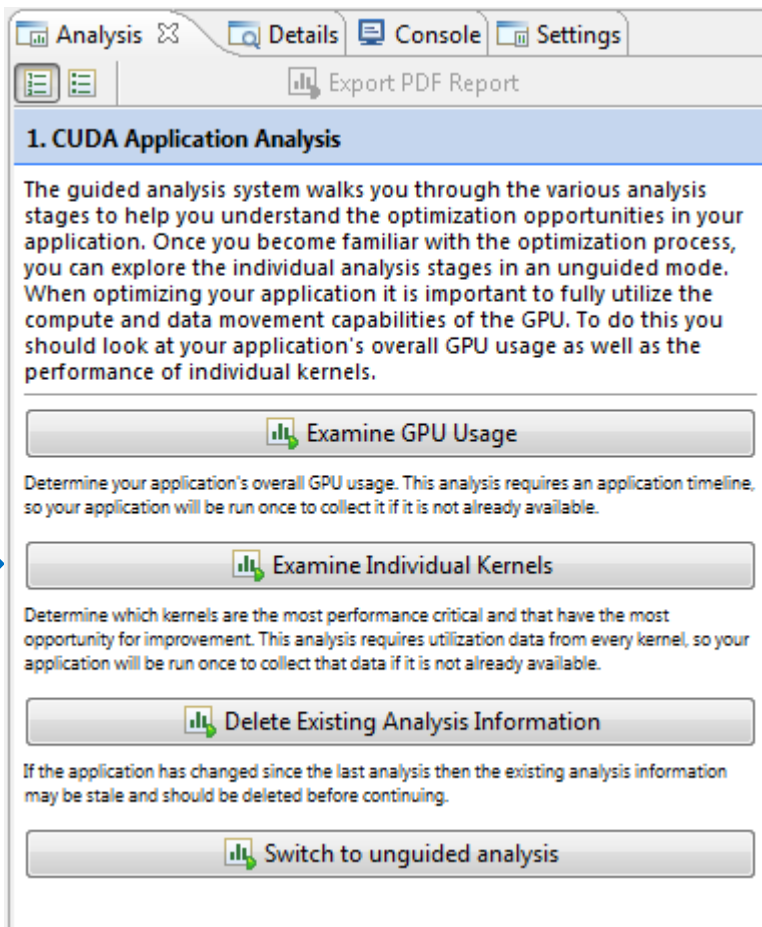


Back to vecadd

More details!

- If we collected our profiling information with -analysis-metrics, or if we're running the application through nvvp directly, more information is available!
- -analysis-metrics may take some time to run
 - Run separately after generating timeline info

This is what we want!



i Kernel Optimization Priorities

The following kernels are ordered by optimization importance based on execution time and achieved occupancy. Optimization of higher ranked kernels (those that appear first in the list) is more likely to improve performance compared to lower ranked kernels.

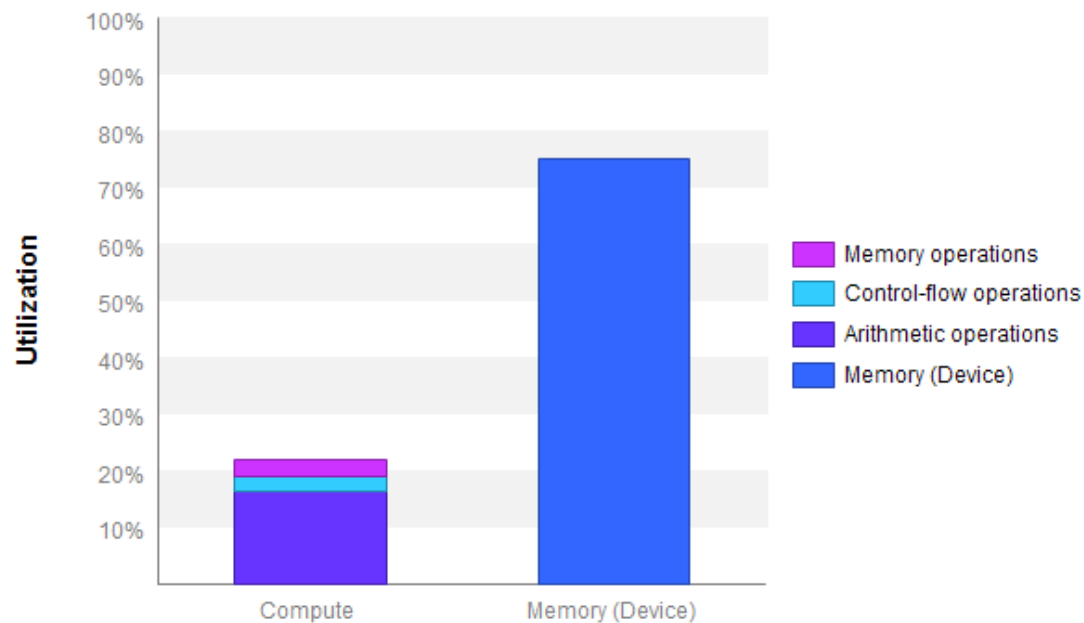
Rank	Description	
100	[1 kernel instances] kernel_vecadd_gpu_	
60	[2 kernel instances] __pgi_dev_cumemset_4	



Perform Kernel Analysis

i Kernel Performance Is Bound By Memory Bandwidth

For device "Tesla K40m" the kernel's compute utilization is significantly lower than its memory utilization. These utilization levels indicate that the performance of the kernel is most likely being limited by the memory system. For this kernel the limiting factor in the memory system is the bandwidth of the Device memory.



- Matches our earlier calculations of ~72% peak bandwidth

Compute, Bandwidth or Latency Bound

- Other analysis options available
- The tool will guide you
 - Becomes more useful the more experience you get

3. Compute, Bandwidth, or Latency Bound

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "kernel_vecadd_gpu_" is most likely limited by memory bandwidth.

Perform Memory Bandwidth Analysis

The most likely bottleneck to performance for this kernel is memory bandwidth so you should first perform memory bandwidth analysis to determine how it is limiting performance.

Perform Compute Analysis

Perform Latency Analysis

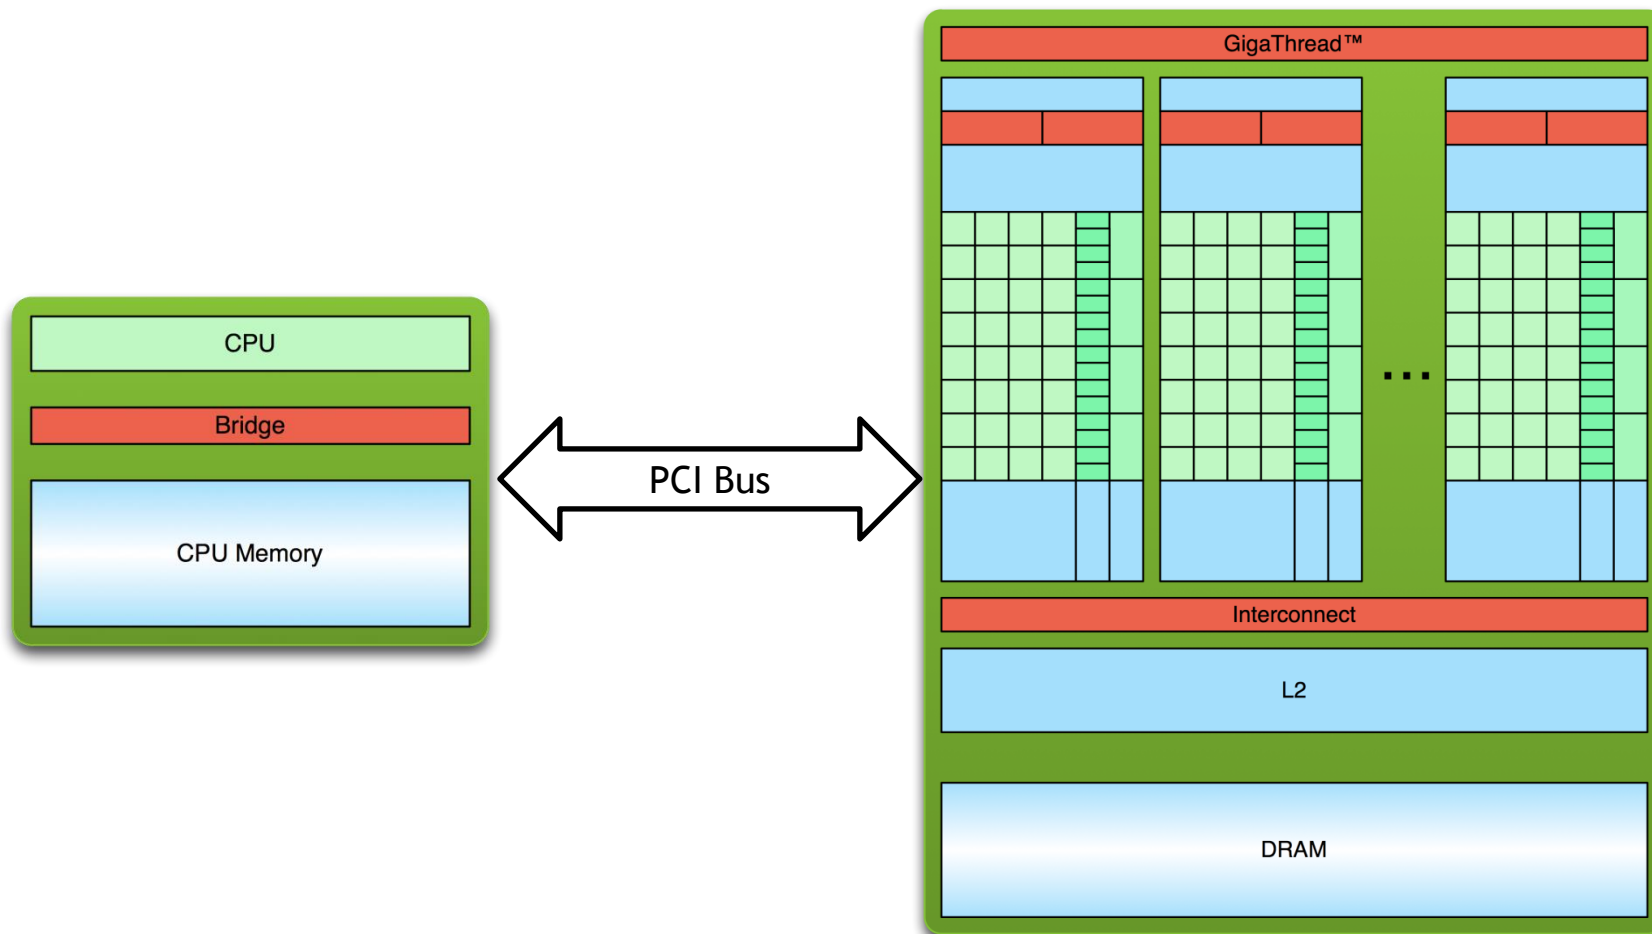
Compute and instruction and memory latency are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

Profiling tips

- The nvprof output is a very good place to start
- The timeline is a good place to go next
- Only dig deep into a kernel if it's taking a significant amount of your time
- Where possible, try to match profiler output with theory
 - For example, if I know I'm moving 1GB, and my kernel takes 10ms, I expect the profiler to report 100GB/s.
 - Discrepancies are likely to mean your application isn't doing what you thought it was

Optimising Data Movement

Heterogeneous Computing



Data movement

What are the costs?

- ▶ Two costs:
 - ▶ Bandwidth
 - ▶ Latency
- ▶ Large copies will be bandwidth bound.
 - ▶ Peak PCI-E bandwidth is 16 GB/s
- ▶ Small copies will be high latency
 - ▶ Many small copies are expensive

Data movement

Several optimisation options

- ▶ Several ways to optimise data movement:
 1. Do less of it!
 2. Fuse small copies into larger ones
 3. Use pinned buffers
 4. Overlap data movement with execution
- ▶ Remember: profile before optimising!

Optimising Data Movement

1. Do less of it!

- Consider moving data movement outside of main loop
- For example, an iterative solver would not want data movement within the iteration loop
- Often, during the process of making an application run on a GPU these additional copies are unavoidable
 - If doing intermediate performance analysis, bear in mind these copies may disappear in the final application
- Consider if data can be passed as an argument (CUDA), or declared as private (OpenACC)

Optimising Data Movement

2. Fuse small copies into larger ones

- Small data transfers are much less efficient
- If the opportunity arises, merging copies can be beneficial
- Can be difficult to accomplish in some applications
 - Additional cost due to packing/unpacking may outweigh the savings

Optimising Data Movement

3. Use pinned buffers

- Page-locked (or pinned) allocations generally have higher bandwidth
- Can harm system performance if overused
- In CUDA:
 - `cudaHostAlloc()` - allocate page-locked memory
 - `cudaHostRegister()` - convert normal allocation to page-locked
 - Pinned attribute in CUDA Fortran
- In OpenACC:
 - (Hopefully) done automatically!

Optimising Data Movement

4. Overlap data movement with execution

- With pinned memory PCI-E transfers can occur at the same time as kernel execution
- In CUDA:
 - `cudaMemcpyAsync`
- In OpenACC:
 - `async` clause on `parallel` or `kernels` directives
- This can greatly reduce the cost of memory copies - almost to none in the right case


Optimising Data Movement

Summary

- There are good (and bad) ways of moving data between GPU and CPU
- For many applications this is not important
 - Data resides on the GPU for the application's lifetime
- Next-gen GPU Pascal will bring hardware improvements - NVLINK

NVLINK

Faster communications between processors

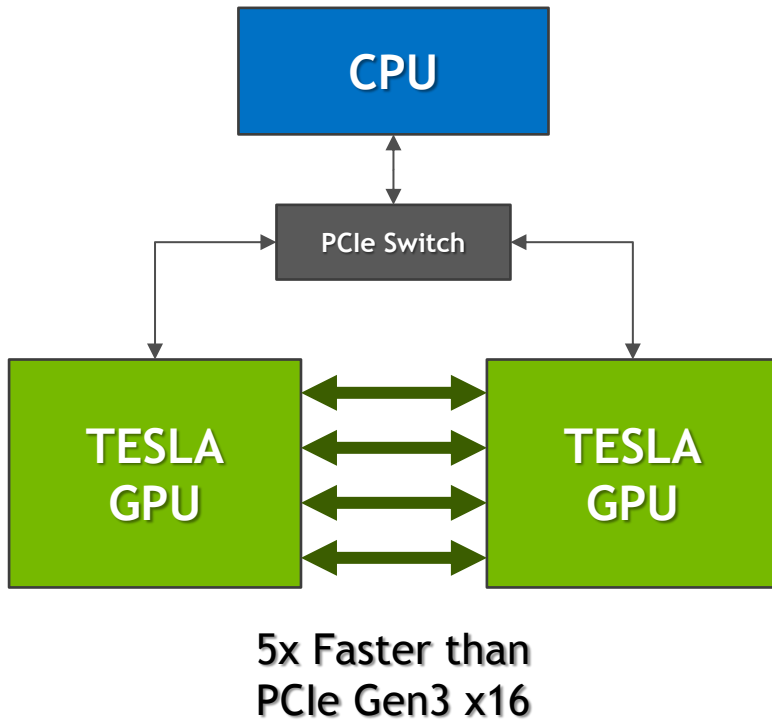


NVLINK

- High speed interconnect
 - Alongside/replacing PCI-E
- 80-200 GB/s
- Improved energy efficiency
- Improved flexibility

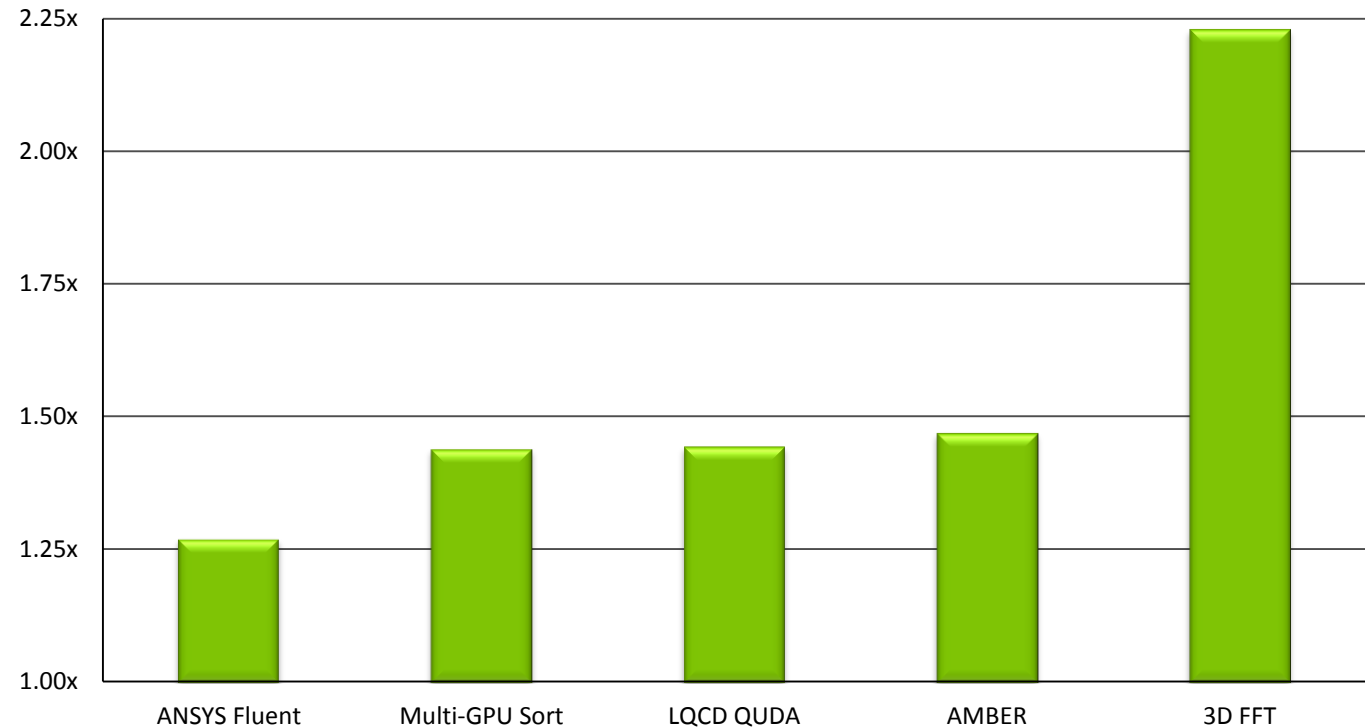
NVLink Unleashes Multi-GPU Performance

GPUs Interconnected with NVLink



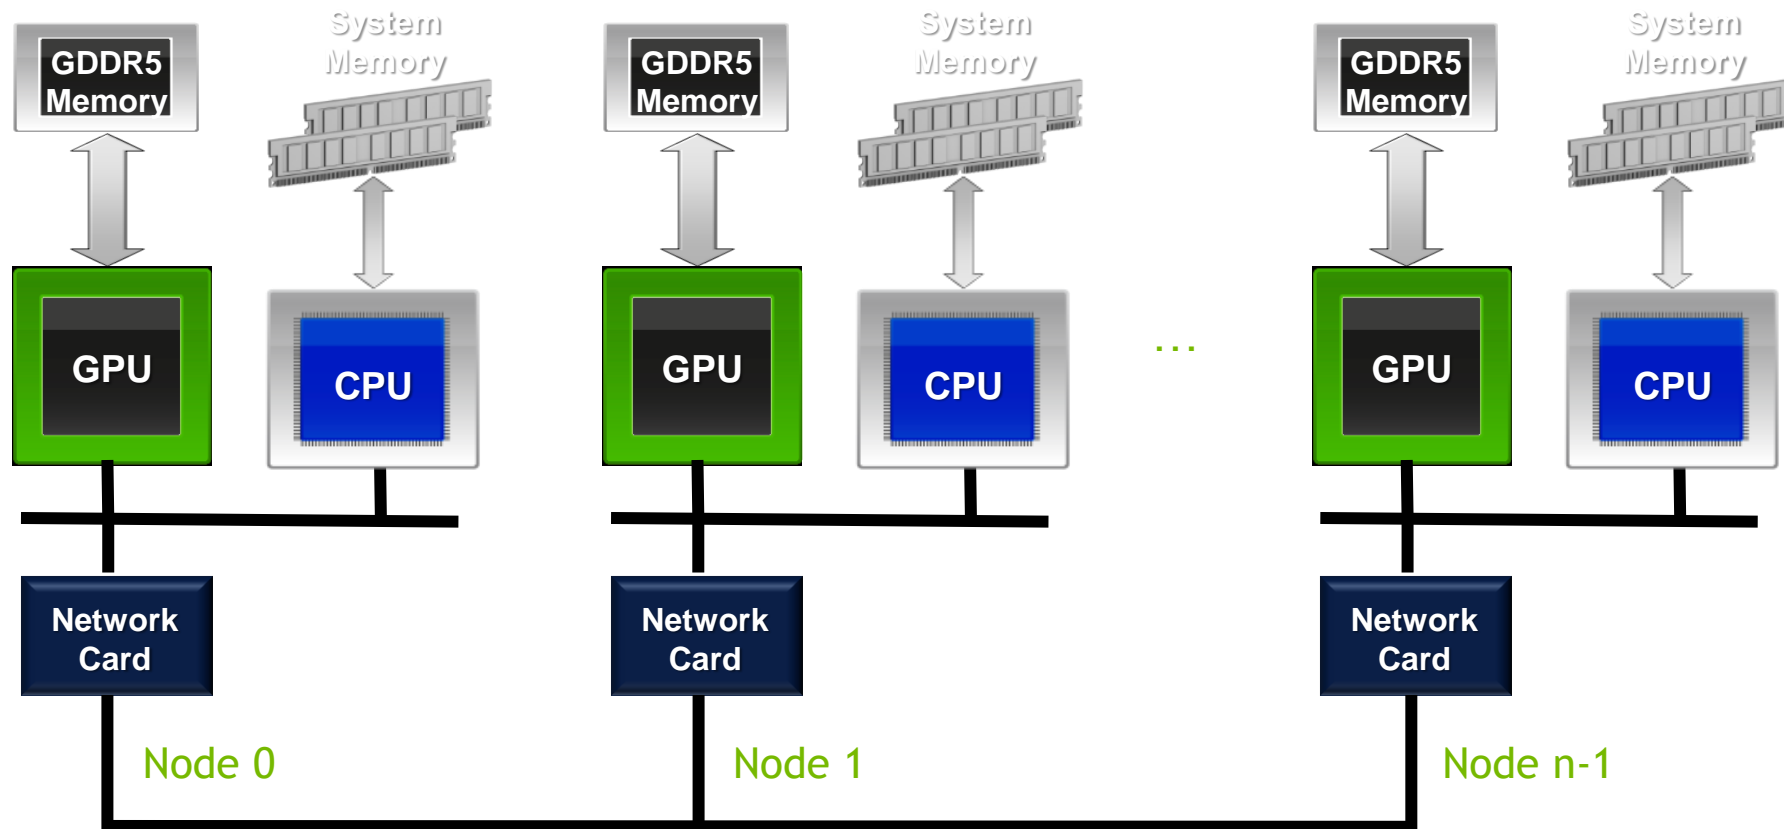
Over 2x Application Performance Speedup When Next-Gen GPUs Connect via NVLink Versus PCIe

Speedup vs
PCIe based Server




Using MPI with GPUs

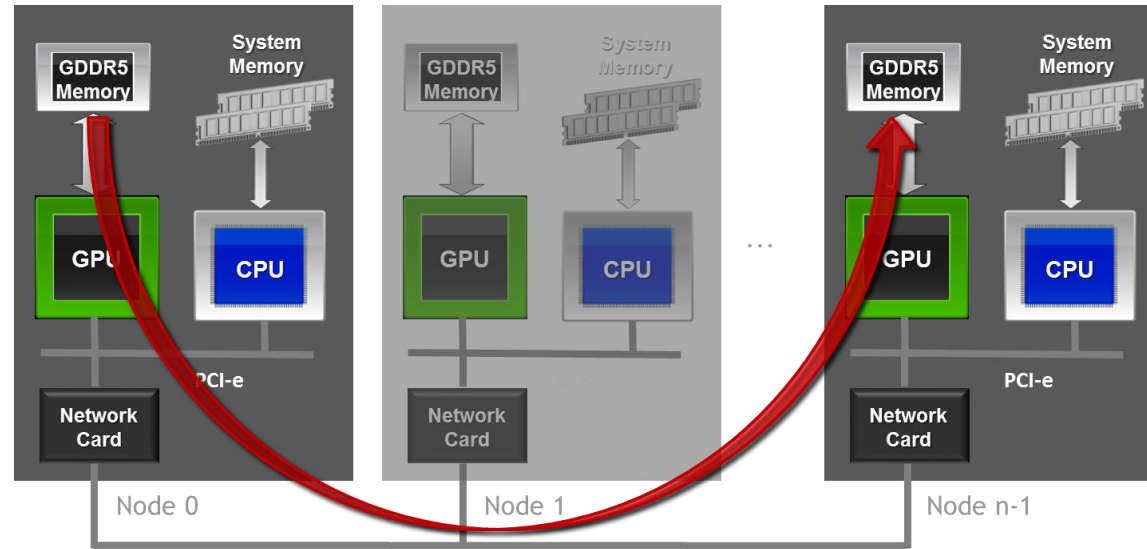
MPI+CUDA



MPI+CUDA



MPI+CUDA



```
//MPI rank 0
```

```
MPI_Send(s_buf_d, size, MPI_CHAR, n-1, tag, MPI_COMM_WORLD);
```

```
//MPI rank n-1
```

```
MPI_Recv(r_buf_d, size, MPI_CHAR, 0, tag, MPI_COMM_WORLD, &stat);
```

Message Passing Interface

MPI

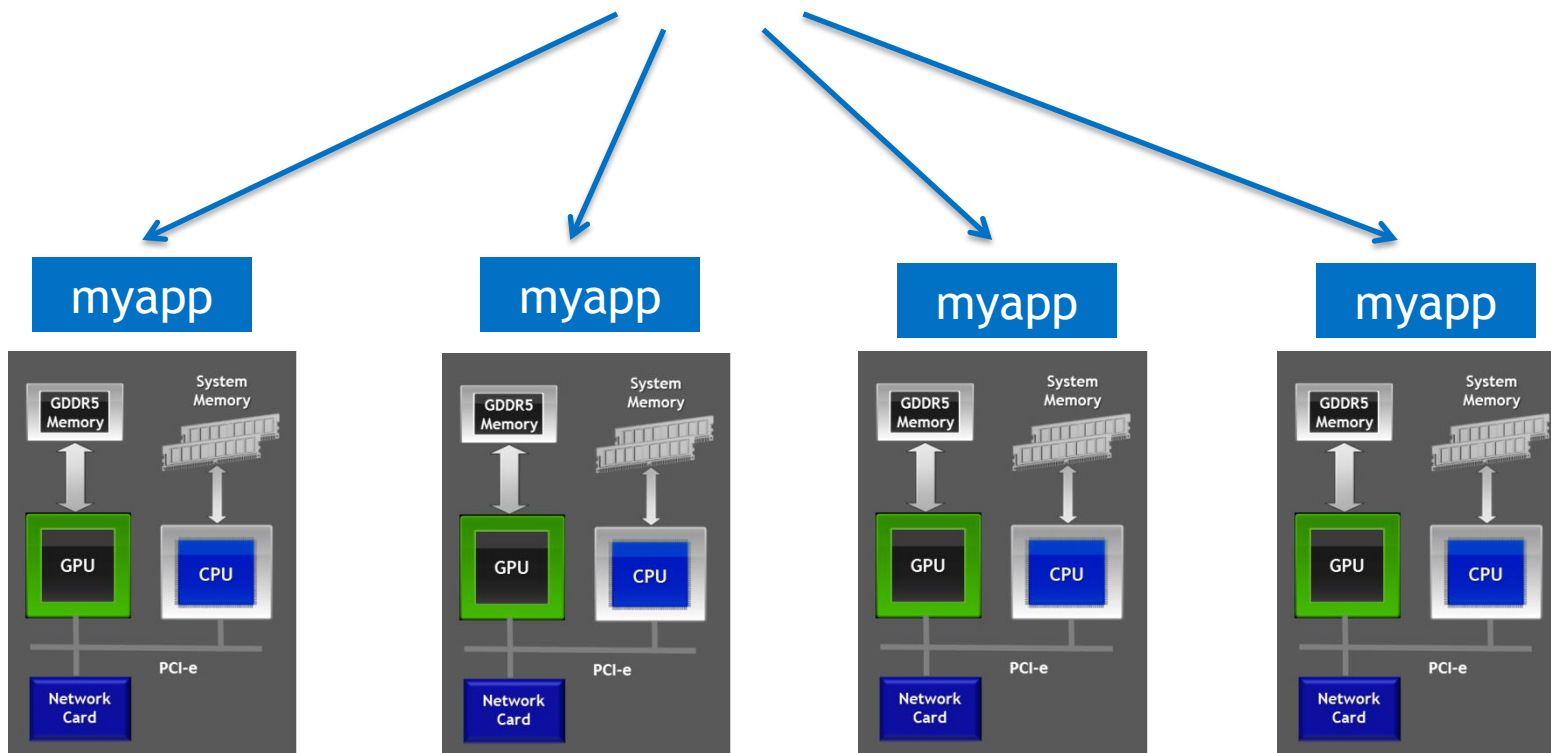
- Standard to exchange data between processes via messages
 - Defines API to exchanges messages
 - Pt. 2 Pt.: e.g. `MPI_Send`, `MPI_Recv`
 - Collectives, e.g. `MPI_Reduce`
- Multiple implementations (open source and commercial)
 - Binding for C/C++, Fortran, Python, ...
 - E.g. MPICH, OpenMPI, MVAPICH, IBM Platform MPI, Cray MPT, ...

MPI

Compiling and launching

```
$ mpicc -o myapp myapp.c
```

```
$ mpirun -np 4 ./myapp <args>
```



Launch MPI + CUDA/OpenACC programs

- Launch one process per GPU

- **MVAPICH:** `MV2_USE_CUDA`

```
$ MV2_USE_CUDA=1 mpirun -np ${np} ./myapp <args>
```

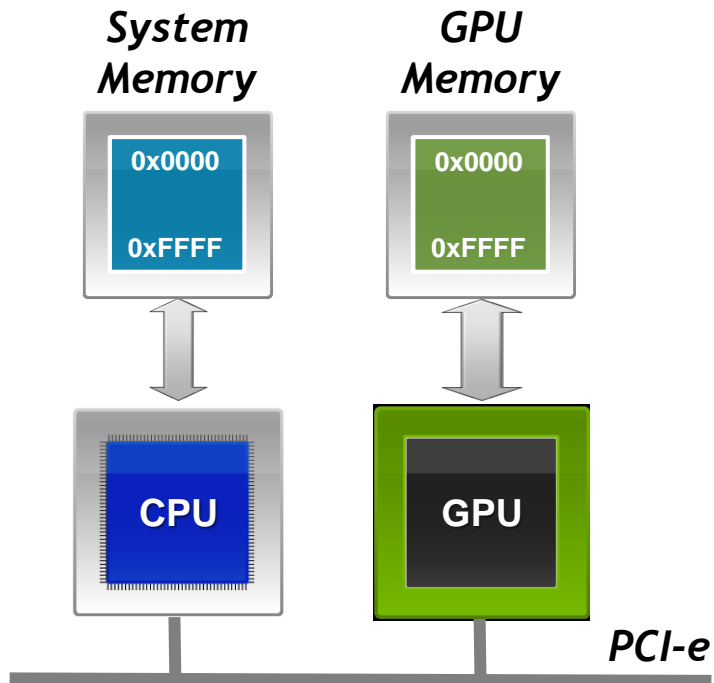
- **Open MPI:** CUDA-aware features are enabled per default
 - **Cray:** `MPICH_RDMA_ENABLED_CUDA`
 - **IBM Platform MPI:** `PMPI_GPU_AWARE`

Unified Virtual Addressing


- One address space for all CPU and GPU memory
 - Determine physical memory location from a pointer value
 - Enable libraries to simplify their interfaces (e.g. MPI and cudaMemcpy)
- Support:
 - 64-bit applications on Linux
 - Windows using TCC mode

Unified Virtual Addressing

No UVA : Separate Address Spaces



UVA : Single Address Space



MPI + CUDA

With UVA CUDA-aware MPI

```
//MPI rank 0  
call MPI_Send(s_buf_d,size,...)  
  
//MPI rank n-1  
call MPI_Recv(r_buf_d,size,...)
```

No UVA and regular MPI

```
//MPI rank 0  
s_buf_h = s_buf_d  
call MPI_Send(s_buf_h,size,...)  
  
//MPI rank n-1  
call MPI_Recv(r_buf_h,size,...)  
r_buf_d = r_buf_h
```

MPI + OpenACC

With UVA CUDA-aware MPI

```
!$acc host_data use_device (s_buf, r_buf)
!MPI rank 0
call MPI_Send(s_buf,size,...)

!MPI rank n-1
call MPI_Recv(r_buf,size,...)
!$acc end host_data
```

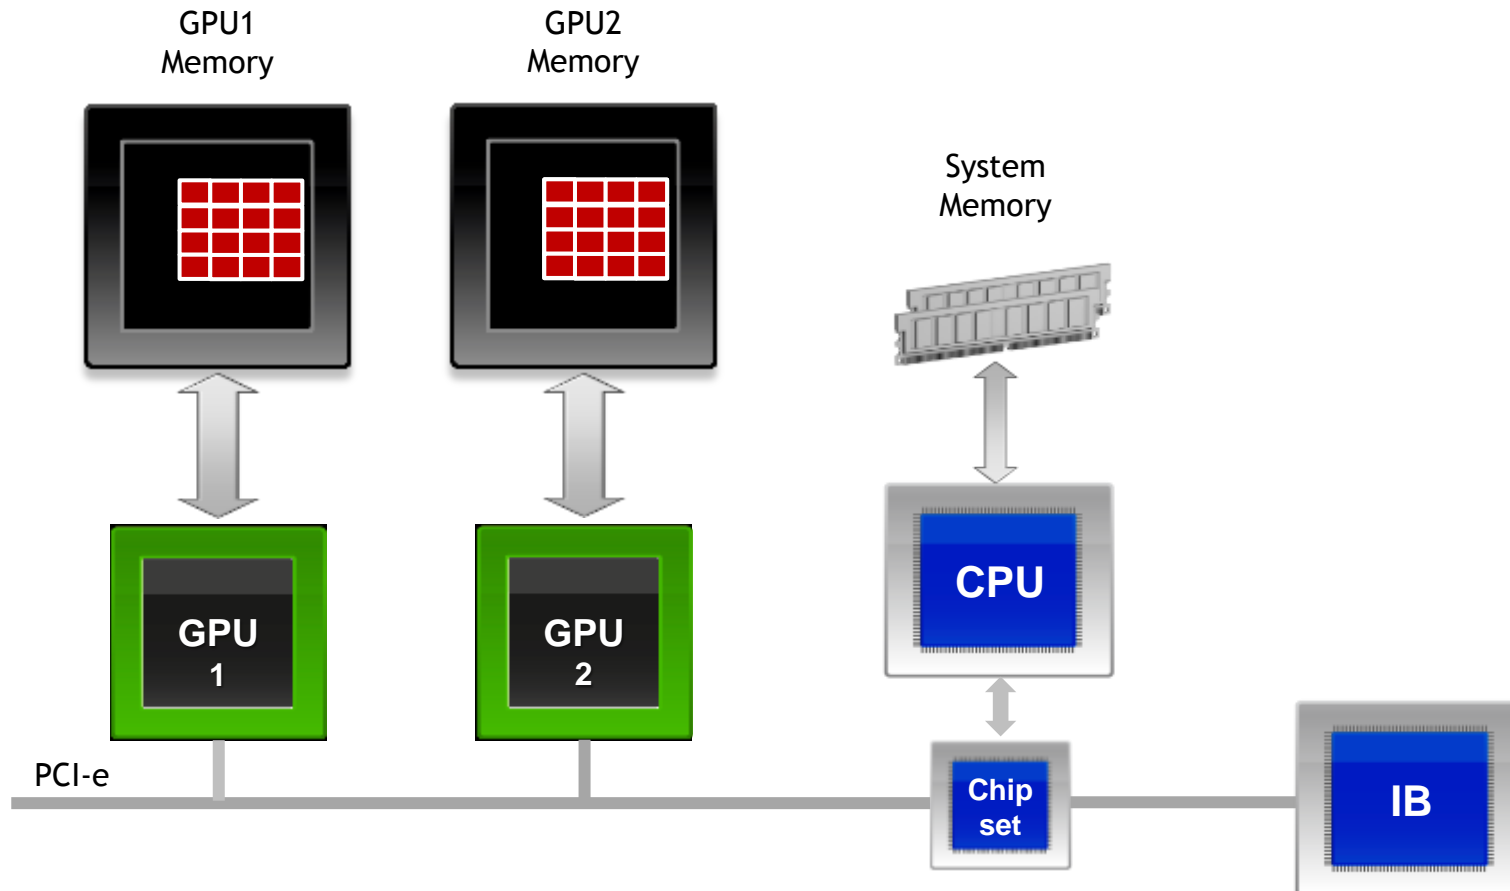
No UVA and regular MPI

```
!$acc update host(s_buf)
!MPI rank 0
call MPI_Send(s_buf,size,...)

!MPI rank n-1
call MPI_Recv(r_buf,size,...)
!$acc update device(r_buf)
```

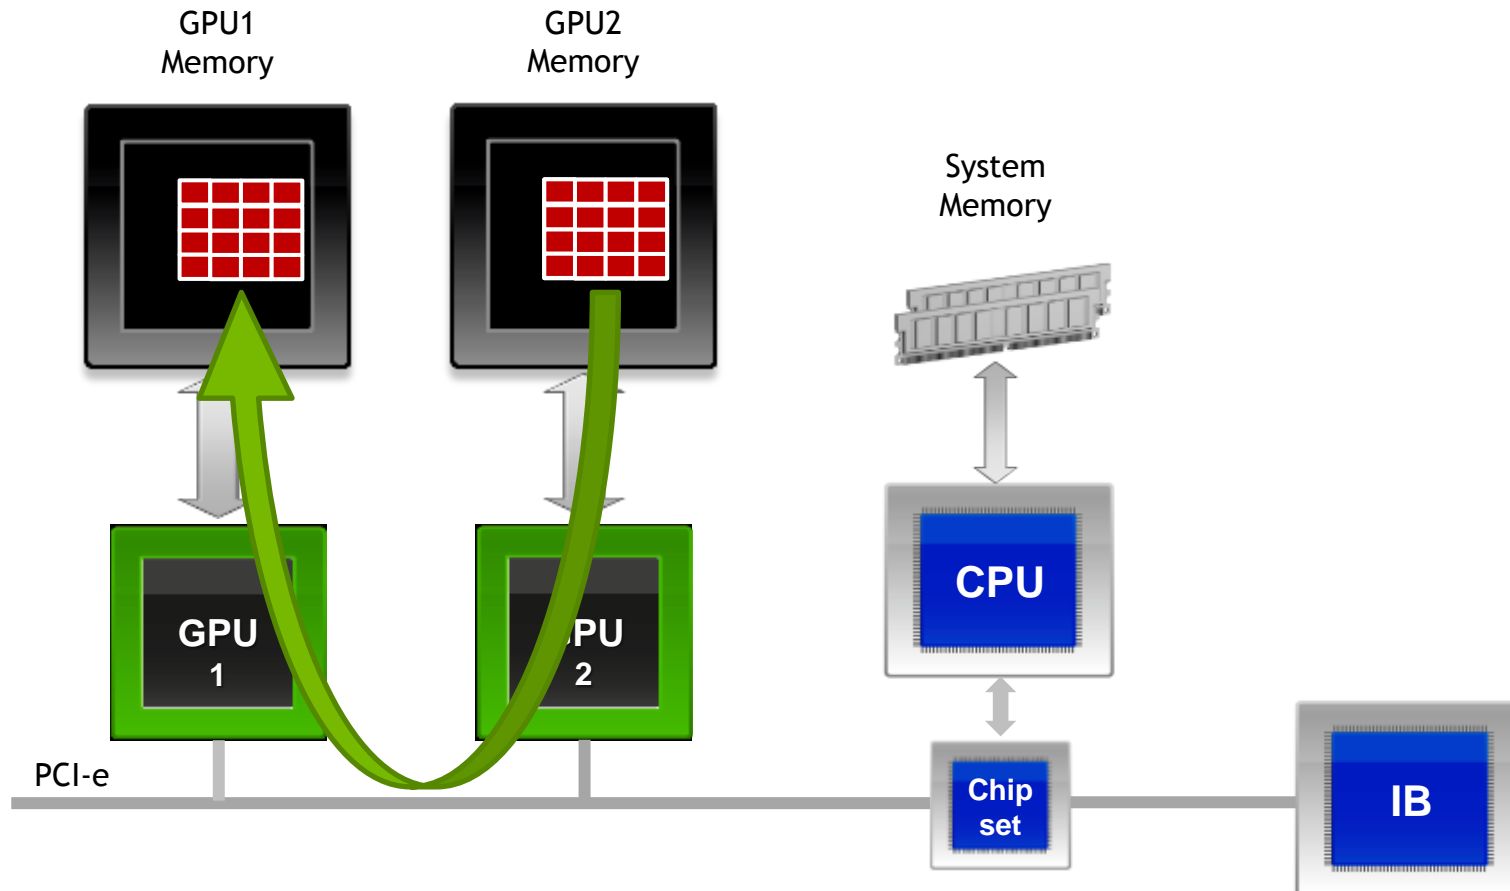
NVIDIA GPUDirect™

Peer to Peer tranfers




NVIDIA GPUDirect™

Peer to Peer tranfers



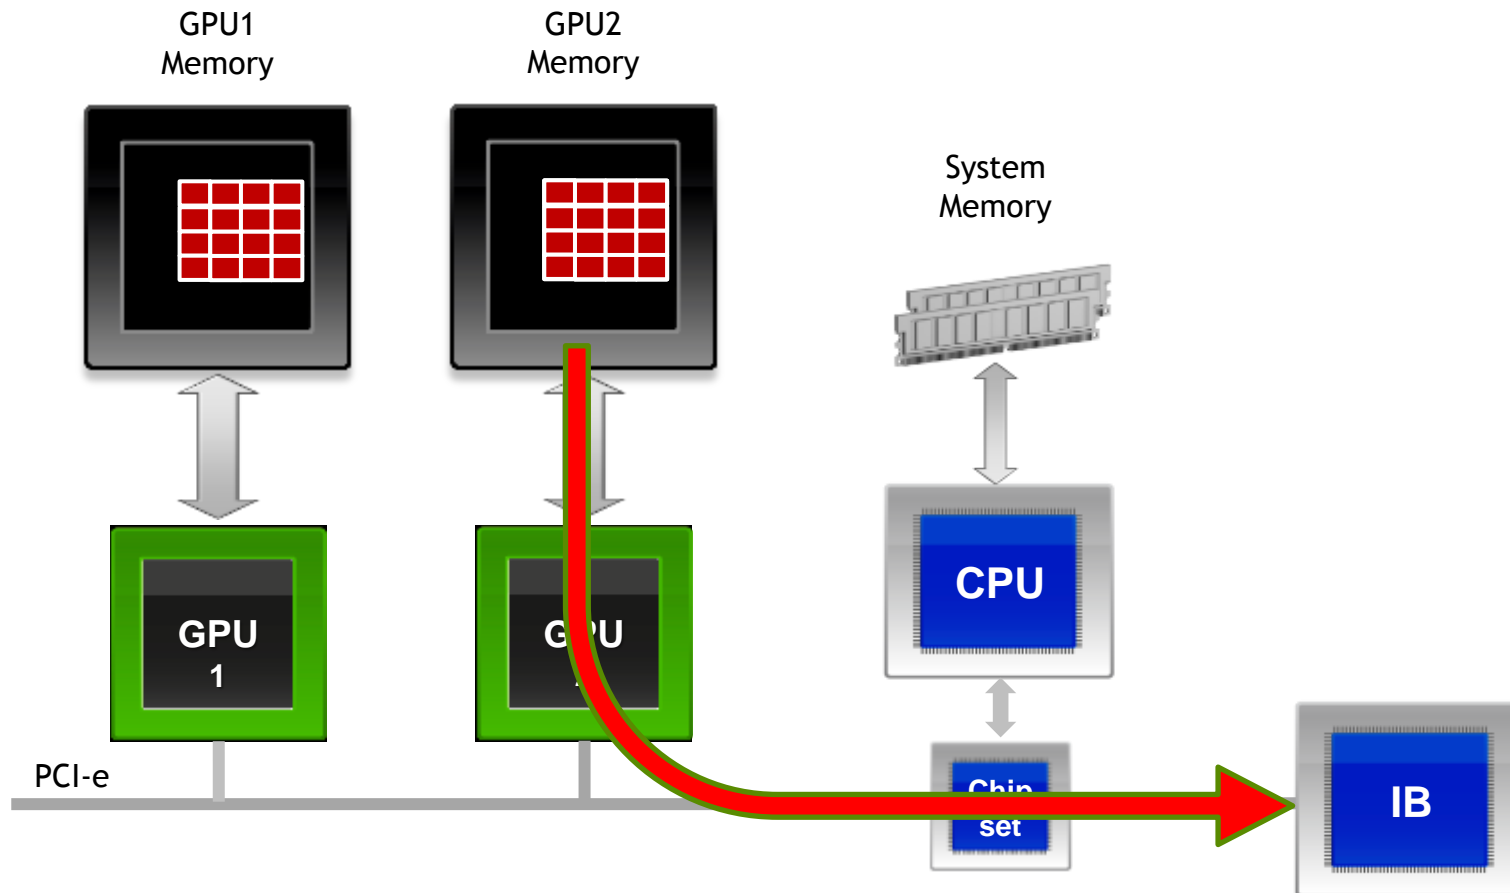
NVIDIA GPUDirect™

Support for RDMA




NVIDIA GPUDirect™

Support for RDMA



Performance Results two Nodes

OpenMPI 1.8.4 MLNX FDR IB (4X) Tesla K40 @ 875GHz



Latency (1 byte) 19.79 us 17.97 us 5.70 us


In this talk

- Profiling GPU applications
- Optimizing Data Movement
- Using MPI with GPUs

Any questions?

Bonus slides - Pascal!

GPU Roadmap




Pascal

- Faster and larger global memory
- Faster communication between processors
- More powerful unified memory
- Mixed precision computing

Stacked Memory

High performance global memory




3D Stacked Memory

- 4x Higher Bandwidth (~1 TB/s)
- 3x Larger Capacity (up to 32GB)
- 4x More Energy Efficient per bit

NVLINK

Faster communications between processors




NVLINK

- High speed interconnect
 - Alongside/replacing PCI-E
- 80-200 GB/s
- Improved energy efficiency
- Improved flexibility



NVLink

High-Speed GPU Interconnect



KEPLER GPU

PASCAL GPU




2014


2016

Unified Memory: Simpler & Faster with NVLink


Traditional Developer View



Developer View With Unified Memory



Developer View With Pascal & NVLink



Mixed precision computing

IEEE 16-bit float support

- Halving precision results in:
 - Half the memory footprint
 - Half the bandwidth required
 - Double the computational throughput
- Obviously comes at an accuracy penalty
 - Application dependent
 - 16/32/64 bit all supported
- Supported in software now